Graded Assignment on Serverless Architecture

Assignment 9: Archive Old Files from S3 to Glacier Using AWS Lambda and Boto3

Objective: Automate the archival of files older than a certain age from an S3 bucket to Amazon Glacier for cost-effective storage.

Task: Automatically move files in an S3 bucket older than 6 months to Glacier storage class.

S3 Setup:

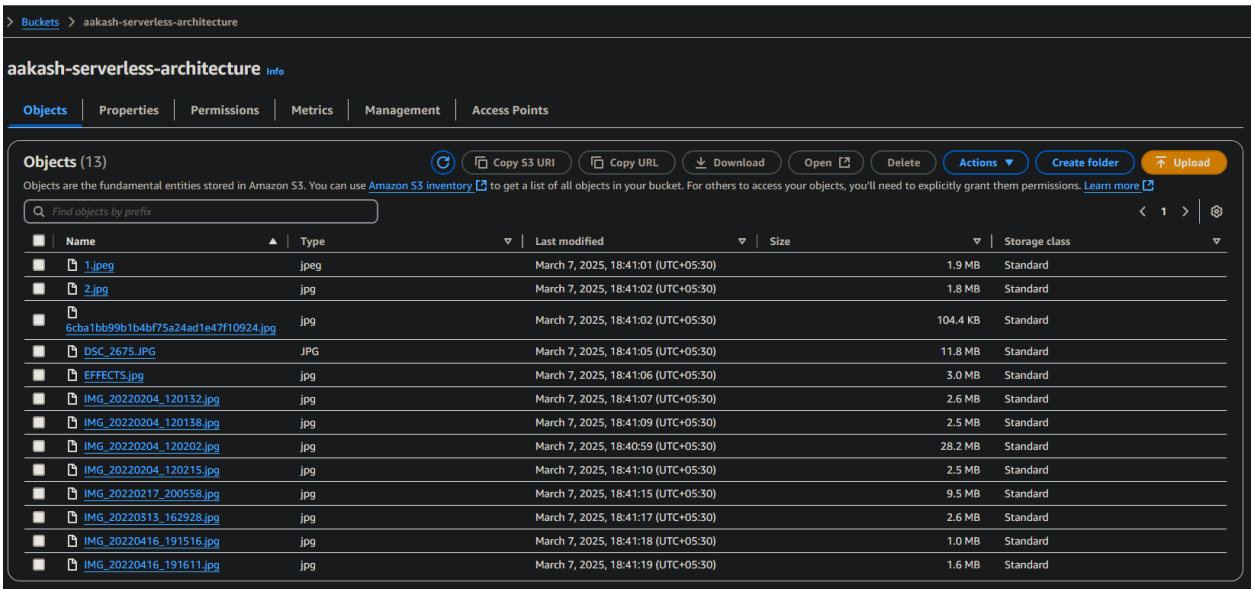

Lambda IAM Role:

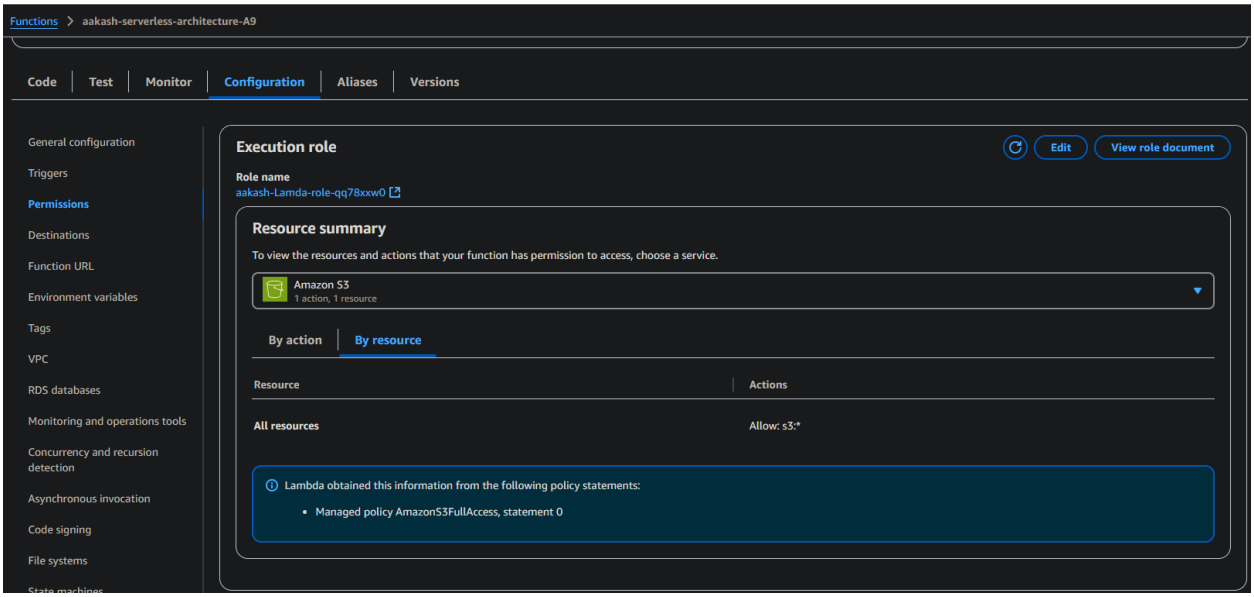

Lambda Function:

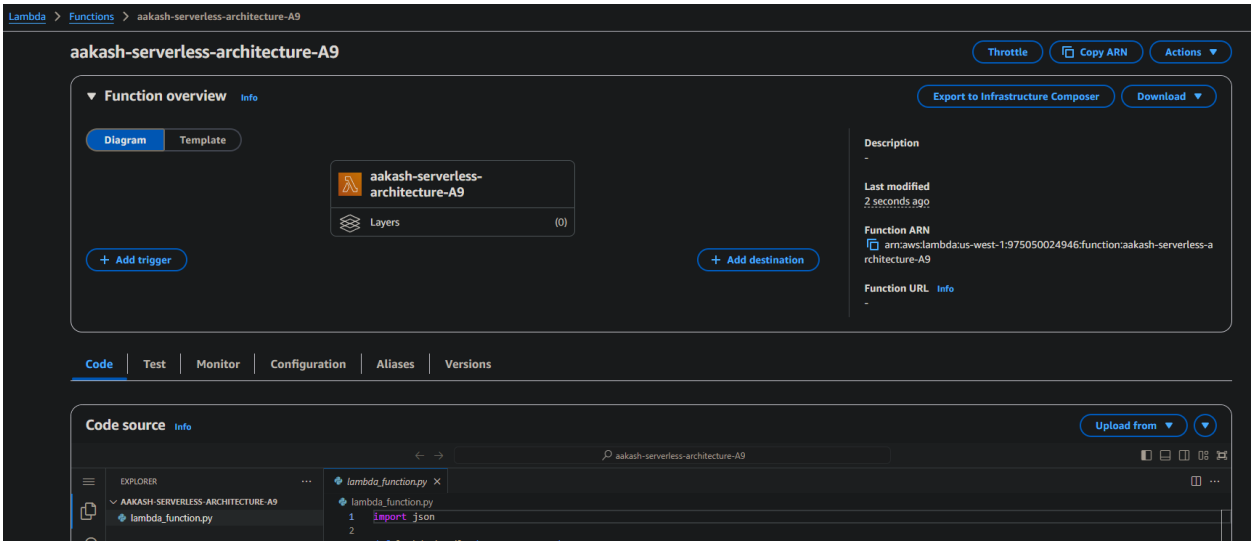

Write the Boto3 Python script:

```
import boto3
import logging
import datetime
from botocore.exceptions import ClientError

# Configure logging
logger = logging.getLogger()
logger.setLevel(logging.INFO)

# S3 client
s3_client = boto3.client('s3')

def lambda_handler(event, context):
    """
    Lambda function to archive S3 objects older than 6 months to Glacier storage
class
    """
    # Set the bucket name - you can also pass this via environment variables
    bucket_name = 'aakash-serverless-architecture'

    # Calculate the date 6 months ago
    today = datetime.datetime.now()
    six_months_ago = today - datetime.timedelta(days=180)

    logger.info(f"Starting archival process for bucket: {bucket_name}")
    logger.info(f"Archiving files older than: {six_months_ago.strftime('%Y-%m-
%d')}")

    try:
        # Get list of objects in the bucket
        response = s3_client.list_objects_v2(Bucket=bucket_name)

        # If no objects are found
        if 'Contents' not in response:
            logger.info(f"No objects found in bucket {bucket_name}")
            return {
```

```
                'statusCode': 200,
                'body': 'No objects found in bucket'
            }

        # Count statistics
        total_objects = 0
        archived_objects = 0

        # Process each object in the bucket
        for obj in response['Contents']:
            total_objects += 1
            key = obj['Key']
            last_modified = obj['LastModified']

            # Convert to datetime object for comparison
            last_modified_date = last_modified.replace(tzinfo=None)

            # Check if the object is older than 6 months
            if last_modified_date < six_months_ago:
                logger.info(f"Archiving object: {key}, Last Modified:
{last_modified_date}")

                # Copy the object with new storage class
                s3_client.copy_object(
                    Bucket=bucket_name,
                    CopySource={'Bucket': bucket_name, 'Key': key},
                    Key=key,
                    StorageClass='GLACIER',
                    MetadataDirective='COPY'
                )

                archived_objects += 1
                logger.info(f"Successfully archived {key} to Glacier")

        # Process additional pages if the response was truncated
        while response.get('IsTruncated', False):
            continuation_token = response.get('NextContinuationToken')
            response = s3_client.list_objects_v2(
                Bucket=bucket_name,
                ContinuationToken=continuation_token
            )

            for obj in response.get('Contents', []):
                total_objects += 1
                key = obj['Key']
```

```
                last_modified = obj['LastModified']

                # Convert to datetime object for comparison
                last_modified_date = last_modified.replace(tzinfo=None)

                # Check if the object is older than 6 months
                if last_modified_date < six_months_ago:
                    logger.info(f"Archiving object: {key}, Last Modified:
{last_modified_date}")

                    # Copy the object with new storage class
                    s3_client.copy_object(
                        Bucket=bucket_name,
                        CopySource={'Bucket': bucket_name, 'Key': key},
                        Key=key,
                        StorageClass='GLACIER',
                        MetadataDirective='COPY'
                    )

                    archived_objects += 1
                    logger.info(f"Successfully archived {key} to Glacier")

        logger.info(f"Archival process completed. Total objects: {total_objects},
Archived: {archived_objects}")

        return {
            'statusCode': 200,
            'body': f"Archival process completed. Total objects: {total_objects},
Archived: {archived_objects}"
        }

    except ClientError as e:
        logger.error(f"Error in archival process: {e}")
        return {
            'statusCode': 500,
            'body': f"Error in archival process: {str(e)}"
        }
```

Testing:

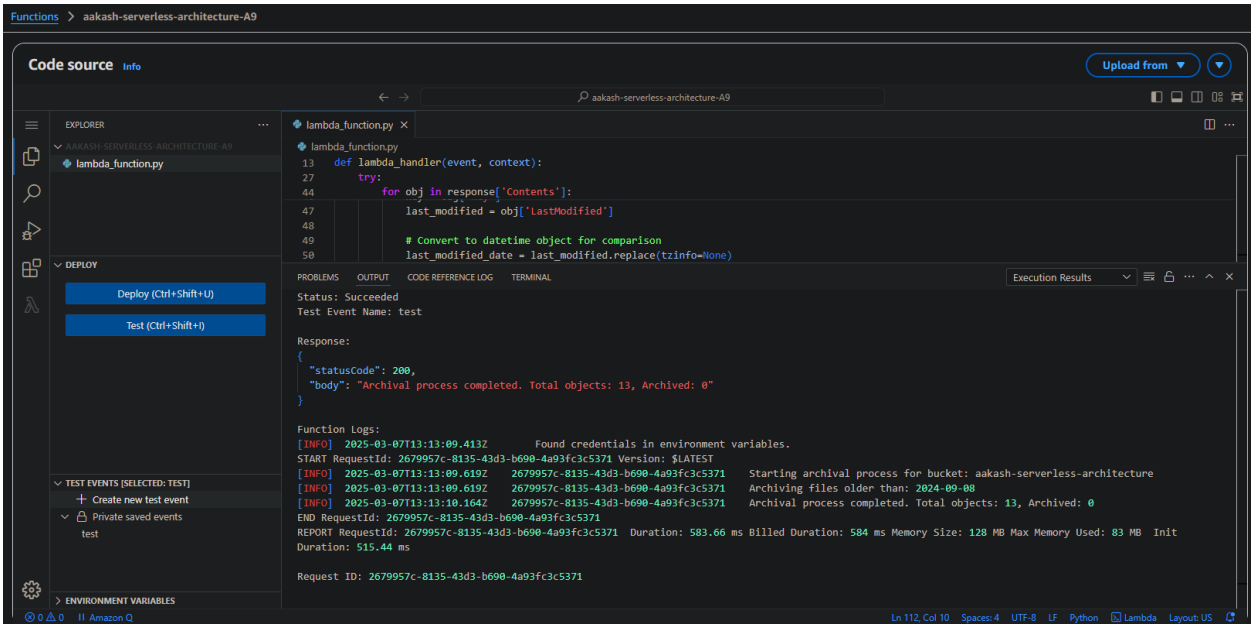
Functions > aakash-serverless-architecture-A9
Code source Info
Upload from
EXPLORER
lambda_function.py
DEPLOY
Deploy (Ctrl+Shift+U)
Test (Ctrl+Shift+I)
TEST EVENTS [SELECTED: TEST]
Create new test event
Private saved events
test
ENVIRONMENT VARIABLES
def lambda_handler(event, context):
try:
for obj in response['Contents']:
last_modified = obj['LastModified']
# Convert to datetime object for comparison
last_modified_date = last_modified.replace(tzinfo=None)
PROBLEMS OUTPUT CODE REFERENCE LOG TERMINAL
Execution Results
Status: Succeeded
Test Event Name: test
Response:
{
"statusCode": 200,
"body": "Archival process completed. Total objects: 13, Archived: 0"
}
Function Logs:
[INFO] 2025-03-07T13:13:09.413Z Found credentials in environment variables.
START RequestId: 2679957c-8135-43d3-b690-4a93fc3c5371 Version: $LATEST
[INFO] 2025-03-07T13:13:09.619Z 2679957c-8135-43d3-b690-4a93fc3c5371 Starting archival process for bucket: aakash-serverless-architecture
[INFO] 2025-03-07T13:13:09.619Z 2679957c-8135-43d3-b690-4a93fc3c5371 Archiving files older than: 2024-09-08
[INFO] 2025-03-07T13:13:10.164Z 2679957c-8135-43d3-b690-4a93fc3c5371 Archival process completed. Total objects: 13, Archived: 0
END RequestId: 2679957c-8135-43d3-b690-4a93fc3c5371
REPORT RequestId: 2679957c-8135-43d3-b690-4a93fc3c5371 Duration: 583.66 ms Billed Duration: 584 ms Memory Size: 128 MB Max Memory Used: 83 MB Init Duration: 515.44 ms
Request ID: 2679957c-8135-43d3-b690-4a93fc3c5371